

Unit-4 Deep Learning

GTU Syllabus:

- Concept of Deep Learning
- Introduction to Neural Networks.
- Types of Deep Learning models.
- Deep learning application

Introduction to Neural Networks:

What is a Neural Network?

Neural networks are machine learning models that mimic the complex functions of the human brain. These models consist of interconnected nodes or neurons that process data, learn patterns, and enable tasks such as pattern recognition and decision-making.

Understanding Neural Networks in Deep Learning

Neural networks are capable of learning and identifying patterns directly from data without pre-defined rules. These networks are built from several key components:

1. **Neurons:** The basic units that receive inputs, each neuron is governed by a threshold and an activation function.
2. **Connections:** Links between neurons that carry information, regulated by weights and biases.
3. **Weights and Biases:** These parameters determine the strength and influence of connections.
4. **Propagation Functions:** Mechanisms that help process and transfer data across layers of neurons.
5. **Learning Rule:** The method that adjusts weights and biases over time to improve accuracy.


Learning in neural networks follows a structured, three-stage process:

1. **Input Computation:** Data is fed into the network.
2. **Output Generation:** Based on the current parameters, the network generates an output.
3. **Iterative Refinement:** The network refines its output by adjusting weights and biases, gradually improving its performance on diverse tasks.

In an adaptive learning environment:

- The neural network is exposed to a simulated scenario or dataset.

- Parameters such as weights and biases are updated in response to new data or conditions.
- With each adjustment, the network's response evolves, allowing it to adapt effectively to different tasks or environments.



The image illustrates the analogy between a biological neuron and an artificial neuron, showing how inputs are received and processed to produce outputs in both systems.

Importance of Neural Networks

Neural networks are pivotal in identifying complex patterns, solving intricate challenges, and adapting to dynamic environments. Their ability to learn from vast amounts of data is transformative, impacting technologies like **natural language processing**, **self-driving vehicles**, and **automated decision-making**.

Neural networks streamline processes, increase efficiency, and support decision-making across various industries. As a backbone of artificial intelligence, they continue to drive innovation, shaping the future of technology.

Evolution of Neural Networks

Neural networks have undergone significant evolution since their inception in the mid-20th century. Here's a concise timeline of the major developments in the field:


- 1940s-1950s:** The concept of neural networks began with McCulloch and Pitts' introduction of the first mathematical model for **artificial neurons**. However, the lack of computational power during that time posed significant challenges to further advancements.
- 1960s-1970s:** Frank Rosenblatt's worked on perceptrons. **Perceptrons** are simple single-layer networks that can solve linearly separable problems, but can not perform complex tasks.
- 1980s:** The development of **backpropagation** by Rumelhart, Hinton, and Williams revolutionized neural networks by enabling the training of multi-layer networks. This

period also saw the rise of connectionism, emphasizing learning through interconnected nodes.

- **1990s:** Neural networks experienced a surge in popularity with applications across image recognition, finance, and more. However, this growth was tempered by a period known as the “**AI winter**,” during which high computational costs and unrealistic expectations dampened progress.
- **2000s:** A resurgence was triggered by the availability of larger datasets, advances in computational power, and innovative network architectures. Deep learning, utilizing multiple layers, proved highly effective across various domains.
- **2010s:** The landscape of machine learning has been dominated by deep learning with **CNNs** (Convolutional Neural Networks) excelling in image classification and **RNNs** (Recurrent Neural Networks) , **LSTMs**, and **GRUs** gaining traction in sequence-based tasks like language modeling and speech recognition.
- **2017: Transformer models**, introduced by Vaswani et al. in “Attention is All You Need,” revolutionized NLP by using a self-attention mechanism for parallel processing, improving efficiency. Models like **BERT**, **GPT**, and **T5** set new benchmarks in machine translation and text generation.

Layers in Neural Network Architecture

1. **Input Layer:** This is where the network receives its input data. Each input neuron in the layer corresponds to a feature in the input data.
2. **Hidden Layers:** These layers perform most of the computational heavy lifting. A neural network can have one or multiple hidden layers. Each layer consists of units (neurons) that transform the inputs into something that the output layer can use.
3. **Output Layer:** The final layer produces the output of the model. The format of these outputs varies depending on the specific task (e.g., classification, regression).



Working of Neural Networks

Forward Propagation

When data is input into the network, it passes through the network in the forward direction, from the input layer through the hidden layers to the output layer. This process is known as forward propagation. Here's what happens during this phase:

1. **Linear Transformation:** Each neuron in a layer receives inputs, which are multiplied by the weights associated with the connections. These products are summed together, and a bias is added to the sum. This can be represented mathematically as:

$z = w_1x_1 + w_2x_2 + \dots + w_nx_n + b$ where w represents the weights, x represents the inputs, and b is the bias.

2. **Activation:** The result of the linear transformation (denoted as z) is then passed through an activation function. The activation function is crucial because it introduces non-linearity into the system, enabling the network to learn more complex patterns. Popular activation functions include ReLU, sigmoid, and tanh.

Note: ReLU, short for rectified linear unit, is a non-linear activation function used for deep neural networks in machine learning. It is also known as the rectifier activation function.

Backpropagation

After forward propagation, the network evaluates its performance using a loss function, which measures the difference between the actual output and the predicted output. The goal of training is to minimize this loss. This is where backpropagation comes into play:

1. **Loss Calculation:** The network calculates the loss, which provides a measure of error in the predictions. The loss function could vary; common choices are mean squared error for regression tasks or cross-entropy loss for classification.
2. **Gradient Calculation:** The network computes the gradients of the loss function with respect to each weight and bias in the network. This involves applying the chain rule of calculus to find out how much each part of the output error can be attributed to each weight and bias.
3. **Weight Update:** Once the gradients are calculated, the weights and biases are updated using an optimization algorithm like stochastic gradient descent (SGD). The weights are adjusted in the opposite direction of the gradient to minimize the loss. The size of the step taken in each update is determined by the learning rate.

Iteration

This process of forward propagation, loss calculation, backpropagation, and weight update is repeated for many iterations over the dataset. Over time, this iterative process reduces the loss, and the network's predictions become more accurate.

Through these steps, neural networks can adapt their parameters to better approximate the relationships in the data, thereby improving their performance on tasks such as classification, regression, or any other predictive modeling.

Example of Email Classification

Let's consider a record of an email dataset:

Email ID	Email Content	Sender	Subject Line	Label
1	"Get free gift cards now!"	spam@example.com	"Exclusive Offer"	1

To classify this email, we will create a feature vector based on the analysis of keywords such as "free," "win," and "offer."

The feature vector of the record can be presented as:

- "free": Present (1)
- "win": Absent (0)
- "offer": Present (1)

Email ID	Email Content	Sender	Subject Line	Feature Vector	Label
1	"Get free gift cards now!"	spam@example.com	"Exclusive Offer"	[1, 0, 1]	1

How Neurons Process Data in a Neural Network

In a **neural network**, input data is passed through multiple layers, including one or more **hidden layers**. Each **neuron** in these hidden layers performs several operations, transforming the input into a usable output.

1. Input Layer: The input layer contains 3 nodes that indicates the presence of each keyword.

2. Hidden Layer

- The input data is passed through one or more hidden layers.
- Each neuron in the hidden layer performs the following operations

Weighted Sum: Each input is multiplied by a corresponding weight assigned to the connection. For example, if the weights from the input layer to the hidden layer neurons are as follows:

- Weights for Neuron H1: [0.5, -0.2, 0.3]
- Weights for Neuron H2: [0.4, 0.1, -0.5]

Calculate Weighted Input:

- For Neuron H1:
 - Calculation = $(1 \times 0.5) + (0 \times -0.2) + (1 \times 0.3) = 0.5 + 0 + 0.3 = 0.8$
- For Neuron H2:
 - Calculation = $(1 \times 0.4) + (0 \times 0.1) + (1 \times -0.5) = 0.4 + 0 - 0.5 = -0.1$

Activation Function: The result is passed through an activation function (e.g., ReLU or sigmoid) to introduce non-linearity.

- For H1, applying ReLU: $\text{ReLU}(0.8) = 0.8$
- For H2, applying ReLU: $\text{ReLU}(-0.1) = 0$


Note: The Rectified Linear Unit (ReLU) activation function is defined by the equation $f(x) = \max(0, x)$. This means that if the input (x) is positive, the output is equal to the input; if the input is negative, the output is zero.

3. Output Layer

- The activated outputs from the hidden layer are passed to the output neuron.
- The output neuron receives the values from the hidden layer neurons and computes the final prediction using weights:
 - Suppose the output weights from hidden layer to output neuron are $[0.7, 0.2]$.
 - Calculation:
 - Input = $(0.8 \times 0.7) + (0 \times 0.2) = 0.56 + 0 = 0.56$
 - **Final Activation:** The output is passed through a sigmoid activation function to obtain a probability:
 - $\sigma(0.56) \approx 0.636$

4. Final Classification

- The output value of approximately **0.636** indicates the probability of the email being spam.
- Since this value is greater than 0.5, the neural network classifies the email as spam (1).



Neural Network for Email Classification Example

Learning of a Neural Network

1. Learning with Supervised Learning

In supervised learning, a neural network learns from labeled input-output pairs provided by a teacher. The network generates outputs based on inputs, and by comparing these outputs to the known desired outputs, an error signal is created. The network iteratively adjusts its parameters to minimize errors until it reaches an acceptable performance level.

2. Learning with Unsupervised Learning

Unsupervised learning involves data without labeled output variables. The primary goal is to understand the underlying structure of the input data (X). Unlike supervised learning, there is no instructor to guide the process. Instead, the focus is on modeling data patterns and relationships, with techniques like clustering and association commonly used.

3. Learning with Reinforcement Learning

Reinforcement learning enables a neural network to learn through interaction with its environment. The network receives feedback in the form of rewards or penalties, guiding it to find an optimal policy or strategy that maximizes cumulative rewards over time. This approach is widely used in applications like gaming and decision-making.

Types of Neural Networks

There are *seven* types of neural networks that can be used.

- **Feedforward Networks:** A feedforward neural network is a simple artificial neural network architecture in which data moves from input to output in a single direction.
- **Singlelayer Perceptron:** A **single-layer perceptron** consists of only one layer of neurons. It takes inputs, applies weights, sums them up, and uses an activation function to produce an output.

- **Multilayer Perceptron (MLP)**: MLP is a type of feedforward neural network with three or more layers, including an input layer, one or more hidden layers, and an output layer. It uses nonlinear activation functions.
- **Convolutional Neural Network (CNN)**: A Convolutional Neural Network (CNN) is a specialized artificial neural network designed for image processing. It employs convolutional layers to automatically learn hierarchical features from input images, enabling effective image recognition and classification.
- **Recurrent Neural Network (RNN)**: An artificial neural network type intended for sequential data processing is called a Recurrent Neural Network (RNN). It is appropriate for applications where contextual dependencies are critical, such as time series prediction and natural language processing, since it makes use of feedback loops, which enable information to survive within the network.
- **Long Short-Term Memory (LSTM)**: LSTM is a type of RNN that is designed to overcome the vanishing gradient problem in training RNNs. It uses memory cells and gates to selectively read, write, and erase information.

Advantages of Neural Networks

Neural networks are widely used in many different applications because of their many benefits:

- **Adaptability**: Neural networks are useful for activities where the link between inputs and outputs is complex or not well defined because they can adapt to new situations and learn from data.
- **Pattern Recognition**: Their proficiency in pattern recognition renders them efficacious in tasks like as audio and image identification, natural language processing, and other intricate data patterns.
- **Parallel Processing**: Because neural networks are capable of parallel processing by nature, they can process numerous jobs at once, which speeds up and improves the efficiency of computations.
- **Non-Linearity**: Neural networks are able to model and comprehend complicated relationships in data by virtue of the non-linear activation functions found in neurons, which overcome the drawbacks of linear models.

Disadvantages of Neural Networks

Neural networks, while powerful, are not without drawbacks and difficulties:

- **Computational Intensity**: Large neural network training can be a laborious and computationally demanding process that demands a lot of computing power.
- **Black box Nature**: As “black box” models, neural networks pose a problem in important applications since it is difficult to understand how they make decisions.
- **Overfitting**: Overfitting is a phenomenon in which neural networks commit training material to memory rather than identifying patterns in the data. Although regularization approaches help to alleviate this, the problem still exists.
- **Need for Large datasets**: For efficient training, neural networks frequently need sizable, labeled datasets; otherwise, their performance may suffer from incomplete or skewed data.

Applications of Neural Networks

Neural networks have numerous applications across various fields:

1. **Image and Video Recognition**: CNNs are extensively used in applications such as facial recognition, autonomous driving, and medical image analysis.
2. **Natural Language Processing (NLP)**: RNNs and transformers power language translation, chatbots, and sentiment analysis.

3. **Finance:** Predicting stock prices, fraud detection, and risk management.
4. **Healthcare:** Neural networks assist in diagnosing diseases, analyzing medical images, and personalizing treatment plans.
5. **Gaming and Autonomous Systems:** Neural networks enable real-time decision-making, enhancing user experience in video games and enabling autonomous systems like self-driving cars.

Different types of activation functions in Neural Networks:

Activation functions introduce non-linearity to the neural network, enabling it to learn complex patterns. Here are the most common types:

1. Linear Activation Function:

- **Formula:** $f(x) = x$
- **Pros:** Simplicity.
- **Cons:** Cannot handle non-linear data; gradient is constant, leading to limited learning capability.

2. Sigmoid (Logistic) Activation Function:

$$f(x) = \frac{1}{1+e^{-x}}$$

- **Formula:**
- **Range:** (0, 1)
- **Pros:** Smooth gradient; output is bounded.
- **Cons:** Vanishing gradient problem; outputs not zero-centered.

3. Tanh (Hyperbolic Tangent) Activation Function:

$$f(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

- **Formula:**
- **Range:** (-1, 1)
- **Pros:** Zero-centered; stronger gradients than Sigmoid.
- **Cons:** Suffers from vanishing gradients.

4. ReLU (Rectified Linear Unit):

- **Formula:** $f(x) = \max(0, x)$
- **Range:** $[0, \infty)$
- **Pros:** Efficient computation; sparse activation.
- **Cons:** Can lead to “dead neurons” if inputs are negative (dying ReLU problem).

5. Leaky ReLU:

- **Formula:** $f(x)=x$ if $x>0$; αx otherwise, where α is a small constant (e.g., 0.01).
- **Range:** $(-\infty, \infty)$
- **Pros:** Mitigates dying ReLU issue; allows small negative values.
- **Cons:** Introduces a small bias in the output.

6. Parametric ReLU (PReLU):


- **Formula:** Similar to Leaky ReLU but with α as a learnable parameter.
- **Pros:** Adaptive to data; mitigates dying ReLU.
- **Cons:** Increased computational cost.

Choosing the Right Activation Function:

- **Classification:** Sigmoid, Softmax.
- **Hidden Layers:** ReLU, Leaky ReLU,
- **Preventing Vanishing Gradient:** ReLU, Leaky ReLU
- **Handling Negative Inputs:** Tanh, Leaky ReLU.

What is Deep Learning?

Deep Learning (DL) is a subset of machine learning that uses neural networks with multiple layers to learn complex patterns in data. It is inspired by the human brain and is particularly effective for tasks involving large datasets and high-dimensional data.



How Deep Learning Works?:

1. **Input Layer:** Receives raw data (e.g., images, text, audio).
2. **Hidden Layers:** Multiple layers where data is processed using weights and biases.
3. **Output Layer:** Provides the final prediction or classification.

Key Components of a Deep Learning Model

1. **Input Layer:**
 - The starting layer that receives the raw data. Each node in this layer represents a feature or input variable.
2. **Hidden Layers:**
 - Layers between the input and output layers where computations occur. They consist of neurons that apply activation functions to the weighted inputs to introduce non-linearity.
3. **Neurons (Nodes):**
 - Basic units in each layer that perform calculations using weights, biases, and activation functions.
4. **Weights:**
 - Parameters that determine the importance of each input feature. They are adjusted during training to minimize the model's error.
5. **Bias:**
 - Additional parameter added to the weighted sum to allow the model to fit the data better, shifting the activation function.
6. **Activation Functions:**
 - Functions that introduce non-linearity to the model, allowing it to learn complex patterns. Common examples: ReLU, Sigmoid, Tanh, and Softmax.
7. **Loss Function (Cost Function):**
 - Measures the discrepancy between the predicted output and the actual output. Examples include Mean Squared Error, Cross-Entropy, and Hinge Loss.
8. **Optimizer:**
 - Algorithm used to update the model's weights to minimize the loss function. Popular optimizers include Gradient Descent, Adam, and RMSprop.
9. **Backpropagation:**
 - A method for updating weights by calculating the gradient of the loss function and propagating it backward through the network.
10. **Epochs and Batch Size:**
 - **Epoch:** One complete pass through the entire training dataset.
 - **Batch Size:** The number of samples processed before updating the model's weights.
11. **Dropout Layer:**
 - A regularization technique that randomly disables a set of neurons during training to prevent overfitting.
12. **Output Layer:**
 - The final layer that produces the model's prediction. The activation function used here depends on the type of problem (e.g., Softmax for classification).

How is Deep Learning different from Machine Learning? (ML Vs DL)

Aspect	Machine Learning (ML)	Deep Learning (DL)
Definition	A subset of AI that enables computers to learn from data using algorithms.	A subset of ML that uses neural networks with multiple layers to learn complex patterns.
Data Dependency	Works well with structured, smaller datasets.	Requires large datasets to perform effectively.
Feature Engineering	Manual feature extraction and selection are necessary.	Automatically extracts features through multiple layers.
Processing Power	Can run on conventional CPUs.	Requires GPUs/TPUs for efficient processing due to high computational demand.
Training Time	Relatively faster due to simpler algorithms.	Significantly longer due to multiple layers and complex architectures.
Interpretability	Easier to interpret and explain.	Often seen as a “black box” due to complex structures.
Examples	Linear Regression, SVM, Decision Trees.	CNNs, RNNs, GANs, Transformers.

What is Perceptron?

Perceptron is a type of neural network that performs binary classification that maps input features to an output decision, usually classifying data into one of two categories, such as 0 or 1.

Perceptron consists of a single layer of input nodes that are fully connected to a layer of output nodes. It is particularly good at learning **linearly separable patterns**. It utilizes a variation of artificial neurons called **Threshold Logic Units (TLU)**, which were first introduced by McCulloch and Walter Pitts in the 1940s. This foundational model has played a crucial role in the development of more advanced neural networks and machine learning algorithms.

Types of Perceptron

1. **Single-Layer Perceptron** is a type of perceptron is limited to learning linearly separable patterns. It is effective for tasks where the data can be divided into distinct categories through a straight line. While powerful in its simplicity, it struggles with more complex problems where the relationship between inputs and outputs is non-linear.
2. **Multi-Layer Perceptron** possess enhanced processing capabilities as they consist of two or more layers, adept at handling more complex patterns and relationships within the data.

Basic Components of Perceptron

A Perceptron is composed of key components that work together to process information and make predictions.

- **Input Features:** The perceptron takes multiple input features, each representing a characteristic of the input data.

- **Weights**: Each input feature is assigned a weight that determines its influence on the output. These weights are adjusted during training to find the optimal values.
- **Summation Function**: The perceptron calculates the weighted sum of its inputs, combining them with their respective weights.
- **Activation Function**: The weighted sum is passed through the **Heaviside step function**, comparing it to a threshold to produce a binary output (0 or 1).
- **Output**: The final output is determined by the activation function, often used for **binary classification** tasks.
- **Bias**: The bias term helps the perceptron make adjustments independent of the input, improving its flexibility in learning.
- **Learning Algorithm**: The perceptron adjusts its weights and bias using a learning algorithm, such as the Perceptron Learning Rule, to minimize prediction errors.

These components enable the perceptron to learn from data and make predictions. While a single perceptron can handle simple binary classification, complex tasks require multiple perceptrons organized into layers, forming a neural network.

How does Perceptron work?

A weight is assigned to each input node of a perceptron, indicating the importance of that input in determining the output. The Perceptron's output is calculated as a weighted sum of the inputs, which is then passed through an activation function to decide whether the Perceptron will fire.


The weighted sum is computed as:

$$z = w_1x_1 + w_2x_2 + \dots + w_nx_n = X^TW$$

The step function compares this weighted sum to a threshold. If the input is larger than the threshold value, the output is 1; otherwise, it's 0. This is the most common activation function used in Perceptrons are represented by the Heaviside step function:

$$h(z) = \begin{cases} 0 & \text{if } z < \text{Threshold} \\ 1 & \text{if } z \geq \text{Threshold} \end{cases}$$

A perceptron consists of a single layer of Threshold Logic Units (TLU), with each TLU fully connected to all input nodes.




Threshold Logic units

Artificial Neural Network:

Artificial Neural Networks contain artificial neurons, which are called **units**. These units are arranged in a series of layers that together constitute the whole **Artificial Neural Network** in a system. A layer can have only a dozen units or millions of units, as this depends on how the complex neural networks will be required to learn the hidden patterns in the dataset. Commonly, **an** Artificial Neural Network has an input layer, an output layer, as well as hidden layers. The input layer receives data from the outside world, which the neural network needs to analyze or learn about. Then, this data passes through one or multiple hidden layers that transform the input into data that is valuable for the output layer. Finally, the output layer provides an output in the form of a response of the **Artificial Neural Networks** to the input data provided.

In the majority of neural networks, units are interconnected from one layer to another. Each of these connections has weights that determine the influence of one unit on another unit. As the data transfers from one unit to another, the neural network learns more and more about the data, which eventually results in an output from the output layer.



Neural Networks Architecture

The structures and operations of human neurons serve as the basis for artificial neural networks. It is also known as neural networks or neural nets. The **input layer** of an artificial neural network is the **first layer**, and it receives input from external sources and releases it to the **hidden layer, which is the second layer**. In the hidden layer, each neuron receives input from the previous layer neurons, computes the weighted sum, and sends it to the neurons in the next layer. These connections are weighted means effects of the inputs from the previous layer are optimized more or less by assigning different-different weights to each input and it is adjusted during the training process by optimizing these weights for improved model performance.

Artificial neurons vs Biological neurons

The concept of artificial neural networks comes from biological neurons found in animal brains. So they share a lot of similarities in structure and function wise.

- **Structure:** The **structure of artificial neural networks** is inspired by biological neurons. A biological neuron has a cell body or soma to process the impulses, dendrites to receive them, and an axon that transfers them to other neurons. The input nodes of artificial neural networks receive input signals, the hidden layer nodes compute these input signals, and the output layer nodes compute the final output by processing the hidden layer's results using activation functions.


Biological Neuron	Artificial Neuron
Dendrite	Inputs
Cell nucleus or Soma	Nodes
Synapses	Weights

Biological Neuron	Artificial Neuron
Axon	Output

- **Synapses:** Synapses are the links between biological neurons that enable the transmission of impulses from dendrites to the cell body. Synapses are the weights that join the one-layer nodes to the next-layer nodes in artificial neurons. The strength of the links is determined by the weight value.
- **Learning:** In biological neurons, learning happens in the cell body nucleus or soma, which has a nucleus that helps to process the impulses. An action potential is produced and travels through the axons if the impulses are powerful enough to reach the threshold. This becomes possible by synaptic plasticity, which represents the ability of synapses to become stronger or weaker over time in reaction to changes in their activity. In artificial neural networks, backpropagation is a technique used for learning, which adjusts the weights between nodes according to the error or differences between predicted and actual outcomes.

Biological Neuron	Artificial Neuron
Synaptic plasticity	Backpropagations

- **Activation:** In biological neurons, activation is the firing rate of the neuron which happens when the impulses are strong enough to reach the threshold. In artificial neural networks, A mathematical function known as an activation function maps the input to the output, and executes activations.



Biological neurons to Artificial neurons

How do Artificial Neural Networks learn?

Artificial neural networks are trained using a training set. For example, suppose you want to teach an ANN to recognize a cat. Then it is shown thousands of different images of cats so that the network can learn to identify a cat. Once the neural network has been trained enough using images of cats, then you need to check if it can identify cat images correctly.

This is done by making the ANN classify the images it is provided by deciding whether they are cat images or not. The output obtained by the ANN is corroborated by a human-provided description of whether the image is a cat image or not. If the ANN identifies incorrectly then it is used to adjust whatever it has learned during training. **Backpropagation** is done by fine-tuning the weights of the connections in ANN units based on the error rate obtained. This process continues until the artificial neural network can correctly recognize a cat in an image with minimal possible error rates.

Backpropagation in Neural Network

What is Backpropagation?

Backpropagation is a supervised learning algorithm used to minimize the error in neural networks by adjusting weights based on the gradient of the loss function. It operates by propagating the error backward from the output layer to the input layer, updating the weights using the gradient descent method.

Steps of Backpropagation:

1. **Forward Propagation:**
 - Input data is passed through the network layer by layer to compute the predicted output.
2. **Calculate Loss:**
 - The loss/error is computed using a loss function (e.g., Mean Squared Error, Cross-Entropy).
3. **Backward Propagation:**
 - The error is propagated backward to each neuron, calculating the gradient of the loss with respect to each weight using the chain rule.
4. **Weight Update:**
 - Weights are adjusted using the gradient and a learning rate η
 - Formula:

$$w_{new} = w_{old} - \eta \times \frac{\partial L}{\partial w}$$

5.Repeat:

- Steps 1–4 are repeated for multiple epochs until the network converges or the desired accuracy is achieved.

Example of Backpropagation:

Consider a simple neural network with one hidden layer and the following structure:

- Input: $x_1=1, x_2=2$
- Weights:

- $w_1=0.4$, $w_2=0.6$ (Input to Hidden)
- $w_3=0.5$, $w_4=0.7$ (Hidden to Output)
- Target Output $y_{target}=0.8$
- Learning Rate $\eta=0.1$
- Activation Function: Sigmoid

Step 1: Forward Propagation

1. Calculate Hidden Layer Input:

$$h = (1 \times 0.4) + (2 \times 0.6) = 1.6$$

2. Apply Sigmoid Activation:

$$h_{activated} = \frac{1}{1 + e^{-1.6}} \approx 0.832$$

3. Calculate Output Layer Input:

$$o = (0.832 \times 0.5) + (0.832 \times 0.7) \approx 1.0$$

4. Apply Sigmoid Activation:

$$o_{activated} = \frac{1}{1 + e^{-1.0}} \approx 0.731$$

Step 2: Calculate Loss

$$L = \frac{1}{2}(y_{target} - o_{activated})^2 = \frac{1}{2}(0.8 - 0.731)^2 \approx 0.0024$$

Step 3: Backward Propagation

1. Calculate Output Layer Gradient:

$$\begin{aligned}\delta_{output} &= (y_{target} - o_{activated}) \times o_{activated} \times (1 - o_{activated}) \\ &= (0.8 - 0.731) \times 0.731 \times (1 - 0.731) \approx 0.0135\end{aligned}$$

2. Calculate Hidden Layer Gradient:

$$\begin{aligned}\delta_{hidden} &= \delta_{output} \times w_3 \times h_{activated} \times (1 - h_{activated}) \\ &= 0.0135 \times 0.5 \times 0.832 \times (1 - 0.832) \approx 0.0015\end{aligned}$$

Step 4: Update Weights:**1. Update Weights from Hidden to Output:**

$$w_3 = 0.5 - 0.1 \times 0.0135 \times 0.832 \approx 0.4989$$

$$w_4 = 0.7 - 0.1 \times 0.0135 \times 0.832 \approx 0.6989$$

2. Update Weights from Input to Hidden:

$$w_1 = 0.4 - 0.1 \times 0.0015 \times 1 \approx 0.39985$$

$$w_2 = 0.6 - 0.1 \times 0.0015 \times 2 \approx 0.5997$$

Step 5: Repeat the Process

- The process is repeated for multiple epochs until the loss is minimized and the model converges.

Deep Learning Models: (OR Types of Neural Networks) [Note: Each model IMP]**Evolution of Neural Architectures**

- The journey of deep learning began with the **perceptron**, a single-layer neural network introduced in the 1950s. While innovative, perceptrons could only solve linearly separable problems, failing at more complex tasks like the XOR problem.
- This limitation led to the development of **Multi-Layer Perceptrons (MLPs)**. It introduced hidden layers and non-linear activation functions. MLPs, trained using **backpropagation**, could model complex, non-linear relationships, marking a significant leap in neural network capabilities.
- This evolution from perceptrons to MLPs laid the groundwork for advanced architectures like CNNs and RNNs, showcasing the power of layered structures in solving real-world problems.

Types of neural networks [Most important question]

- 1. Feedforward neural networks (FNNs)** are the simplest type of ANN, where data flows in one direction from input to output. It is used for basic tasks like classification.
- 2. Convolutional Neural Networks (CNNs)** are specialized for processing grid-like data, such as images. CNNs use convolutional layers to detect spatial hierarchies, making them ideal for computer vision tasks.
- 3. Recurrent Neural Networks (RNNs)** are able to process sequential data, such as time series and natural language. RNNs have loops to retain information over time, enabling applications like language modeling and speech recognition. Variants like LSTMs and GRUs address vanishing gradient issues.

4. Generative Adversarial Networks (GANs) consist of two networks—a generator and a discriminator—that compete to create realistic data. GANs are widely used for image generation, style transfer, and data augmentation.

5. Autoencoders are unsupervised networks that learn efficient data encodings. They compress input data into a latent representation and reconstruct it, useful for dimensionality reduction and anomaly detection.

6. Transformer Networks has revolutionized NLP with self-attention mechanisms. Transformers excel at tasks like translation, text generation, and sentiment analysis, powering models like GPT and BERT.

Feedforward Neural Network:


Feedforward Neural Network (FNN) is a type of artificial neural network in which information flows in a single direction—from the input layer through hidden layers to the output layer—without loops or feedback. It is mainly used for pattern recognition tasks like image and speech classification.

Structure of a Feedforward Neural Network

Feedforward Neural Networks have a structured layered design where data flows sequentially through each layer.

1. **Input Layer:** The input layer consists of neurons that receive the input data. Each neuron in the input layer represents a feature of the input data.
2. **Hidden Layers:** One or more hidden layers are placed between the input and output layers. These layers are responsible for learning the complex patterns in the data. Each neuron in a hidden layer applies a weighted sum of inputs followed by a non-linear activation function.
3. **Output Layer:** The output layer provides the final output of the network. The number of neurons in this layer corresponds to the number of classes in a classification problem or the number of outputs in a regression problem.

Each connection between neurons in these layers has an associated weight that is adjusted during the training process to minimize the error in predictions.



Feed Forward Neural Network

Activation Functions

Activation functions introduce non-linearity into the network enabling it to learn and model complex data patterns.


Common activation functions include:

- **Sigmoid:** $\sigma(x) = \sigma(x) = \frac{1}{1+e^{-x}}$
- **Tanh:** $\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$
- **ReLU:** $\text{ReLU}(x) = \max(0, x)$

Training a Feedforward Neural Network

Training a Feedforward Neural Network involves adjusting the weights of the neurons to minimize the error between the predicted output and the actual output. This process is typically performed using backpropagation and gradient descent.

1. **Forward Propagation:** During forward propagation the input data passes through the network and the output is calculated.
2. **Loss Calculation:** The loss (or error) is calculated using a loss function such as Mean Squared Error (MSE) for regression tasks or Cross-Entropy Loss for classification tasks.
3. **Backpropagation:** In backpropagation the error is propagated back through the network to update the weights. The gradient of the loss function with respect to each weight is calculated and the weights are adjusted using gradient descent.



Forward Propagation

Gradient Descent

Gradient Descent is an optimization algorithm used to minimize the loss function by iteratively updating the weights in the direction of the negative gradient. Common variants of gradient descent include:

- **Batch Gradient Descent:** Updates weights after computing the gradient over the entire dataset.
- **Stochastic Gradient Descent (SGD):** Updates weights for each training example individually.
- **Mini-batch Gradient Descent:** It Updates weights after computing the gradient over a small batch of training examples.

Evaluation of Feedforward neural network

Evaluating the performance of the trained model involves several metrics:


- **Accuracy:** The proportion of correctly classified instances out of the total instances.
- **Precision:** The ratio of true positive predictions to the total predicted positives.
- **Recall:** The ratio of true positive predictions to the actual positives.
- **F1 Score:** The harmonic mean of precision and recall, providing a balance between the two.
- **Confusion Matrix:** A table used to describe the performance of a classification model, showing the true positives, true negatives, false positives, and false negatives.

Convolutional Neural Networks (CNNs)

A **Convolutional Neural Network (CNN)** is a type of deep learning model designed for processing structured grid data, such as images. CNNs effectively extract spatial and temporal features using convolutional and pooling layers.

Convolutional Neural Network (CNN) is an advanced version of artificial neural networks (ANNs), primarily designed to extract features from grid-like matrix datasets. This is particularly useful for visual datasets such as images or videos, where data patterns play a crucial role. CNNs are widely used in computer vision applications due to their effectiveness in processing visual data.

CNNs consist of multiple layers like the input layer, Convolutional layer, pooling layer, and fully connected layers.



Simple CNN architecture

Components of a CNN:

1. **Input Layer:**
 - Receives the raw pixel values of the image (e.g., 32x32x3 for a color image).
2. **Convolutional Layer (Conv Layer):**
 - Applies filters (kernels) to the input data to extract features (edges, textures, etc.).
 - Output: Feature maps.

Formula:

$$Z=(X*W)+b$$

where X is the input, W is the filter, and b is the bias.


3. **Activation Function:**
 - Introduces non-linearity using functions like ReLU, Tanh, or Sigmoid.
4. **Pooling Layer:**

- Reduces spatial dimensions (e.g., Max Pooling or Average Pooling) to minimize computational complexity and prevent overfitting.
- 5. **Fully Connected (FC) Layer:**
 - Connects flattened data to neurons to learn non-linear combinations of features.
 - Often uses activation functions like Softmax for classification.
- 6. **Output Layer:**
 - Provides the final classification or regression output.

How Convolutional Layers Works?

Convolution Neural Networks are neural networks that share their parameters.

Imagine you have an image. It can be represented as a cuboid having its length, width (dimension of the image), and height (i.e the channel as images generally have red, green, and blue channels).



Now imagine taking a small patch of this image and running a small neural network, called a filter or kernel on it, with say, K outputs and representing them vertically.

Now slide that neural network across the whole image, as a result, we will get another image with different widths, heights, and depths. Instead of just R, G, and B channels now we have more channels but lesser width and height. This operation is called Convolution. If the patch size is the same as that of the image it will be a regular neural network. Because of this small patch, we have fewer weights.

Example of CNN: Image Classification

Problem: Classify handwritten digits (0-9) using the **MNIST dataset**.

Steps:

1. **Input Data:** MNIST dataset consisting of 28x28 grayscale images.
2. **Network Architecture:**
 - Input Layer: 28x28x1 (1 channel for grayscale)
 - Conv Layer 1: 32 filters of 3x3, ReLU activation
 - Max Pooling: 2x2

- Conv Layer 2: 64 filters of 3x3, ReLU activation
- Max Pooling: 2x2
- Flatten Layer
- Dense Layer: 128 neurons, ReLU activation
- Output Layer: 10 neurons, Softmax activation

3. Training:

- Loss function: Categorical Cross-Entropy
- Metrics: Accuracy


Advantages of CNNs:

- Automatic feature extraction
- Reduced number of parameters through shared weights
- Spatial hierarchies of features
- Effective in visual data processing

Recurrent Neural Networks (RNNs) :

Recurrent Neural Networks (RNNs) work a bit different from regular neural networks. In neural network the information flows in one direction from input to output. However in RNN information is fed back into the system after each step. Think of it like reading a sentence, when you're trying to predict the next word you don't just look at the current word but also need to remember the words that came before to make accurate guess.

RNNs allow the network to "remember" past information by feeding the output from one step into next step. This helps the network understand the context of what has already happened and make better predictions based on that. For example when predicting the next word in a sentence the RNN uses the previous words to help decide what word is most likely to come next.



Recurrent Neural Network


This image showcases the basic architecture of RNN and the **feedback loop mechanism where the output is passed back as input for the next time step.**

How RNN Differs from Feedforward Neural Networks?

Feedforward Neural Networks (FNNs) process data in one direction from input to output without retaining information from previous inputs. This makes them suitable for tasks with

independent inputs like image classification. However FNNs struggle with sequential data since they lack memory.


Recurrent Neural Networks (RNNs) solve this by **incorporating loops that allow information from previous steps to be fed back into the network**. This feedback enables RNNs to remember prior inputs making them ideal for tasks where context is important.



Key Components of RNNs

1. Recurrent Neurons

The fundamental processing unit in RNN is a **Recurrent Unit**. Recurrent units hold a hidden state that maintains information about previous inputs in a sequence. Recurrent units can "remember" information from prior steps by feeding back their hidden state, allowing them to capture dependencies across time.




Recurrent Neuron

2. RNN Unfolding

RNN unfolding or unrolling is the process of expanding the recurrent structure over time steps. During unfolding each step of the sequence is represented as a separate layer in a series illustrating how information flows across each time step.

This unrolling enables **backpropagation through time (BPTT)** a learning process where errors are propagated across time steps to adjust the network's weights enhancing the RNN's ability to learn dependencies within sequential data.



RNN Unfolding

Recurrent Neural Network Architecture

RNNs share similarities in input and output structures with other deep learning architectures but differ significantly in how information flows from input to output. Unlike traditional deep neural networks, where each dense layer has distinct weight matrices, RNNs use shared weights across time steps, allowing them to remember information over sequences.

In RNNs, the hidden state h_i is calculated for every input X_i to retain sequential dependencies. The computations follow these core formulas:

1. Hidden State Calculation:

$$h = \sigma(U \cdot X + W \cdot h_{t-1} + B)$$

Here, h represents the current hidden state, U and W are weight matrices, and B is the bias.

2. Output Calculation:

$$Y = O(V \cdot h + C)$$


The output Y is calculated by applying O , an activation function, to the weighted hidden state, where V and C represent weights and bias.

3. Overall Function:

$$Y = f(X, h, W, U, V, B, C)$$

This function defines the entire RNN operation, where the state matrix S holds each element s_i representing the network's state at each time step i .

RECURRENT NEURAL NETWORKS



Recurrent Neural Architecture

How does RNN work?

At each time step RNNs process units with a fixed activation function. These units have an internal hidden state that acts as memory that retains information from previous time steps. This memory allows the network to store past knowledge and adapt based on new inputs.

Updating the Hidden State in RNNs

The current hidden state h_t depends on the previous state h_{t-1} and the current input x_t , and is calculated using the following relations:

1. State Update:

$$h_t = f(h_{t-1}, x_t)$$

where:

- h_t is the current state
- h_{t-1} is the previous state
- x_t is the input at the current time step

2. Activation Function Application:

$$h_t = \tanh(Whh \cdot h_{t-1} + Wxh \cdot x_t)$$

Here, Whh is the weight matrix for the recurrent neuron, and Wxh is the weight matrix for the input neuron.

3. Output Calculation:

$$y_t = Why \cdot h_t$$


where y_t is the output and Why is the weight at the output layer.

*These parameters are updated using backpropagation. However, since RNN works on sequential data here we use an updated backpropagation which is known as **backpropagation through time**.*

Backpropagation Through Time (BPTT) in RNNs

Since RNNs process sequential data **Backpropagation Through Time (BPTT)** is used to update the network's parameters. The loss function $L(\theta)$ depends on the final hidden state h_3 and each hidden state relies on preceding ones forming a sequential dependency chain:

h_3 depends on h_2 , h_2 depends on h_1 , ..., h_1 depends on h_0
 h_3 depends on h_1 , ..., h_1 depends on h_0 .



In BPTT, gradients are backpropagated through each time step. This is essential for updating network parameters based on temporal dependencies.

1. **Simplified Gradient Calculation:**

$$\frac{\partial L(\theta)}{\partial W} = \frac{\partial L(\theta)}{\partial h_3} \cdot \frac{\partial h_3}{\partial W}$$

2. **Handling Dependencies in Layers:**

Each hidden state is updated based on its dependencies:

$$h_3 = \sigma(W \cdot h_2 + b)$$

The gradient is then calculated for each state, considering dependencies from previous hidden states.

3. **Gradient Calculation with Explicit and Implicit Parts:** The gradient is broken down into explicit and implicit parts summing up the indirect paths from each hidden state to the weights.

$$\frac{\partial h_3}{\partial W} = \frac{\partial h_3^+}{\partial W} + \frac{\partial h_3}{\partial h_2} \cdot \frac{\partial h_2^+}{\partial W}$$

4. **Final Gradient Expression:**

The final derivative of the loss function with respect to the weight matrix W is computed:

$$\frac{\partial L(\theta)}{\partial W} = \frac{\partial L(\theta)}{\partial h_3} \cdot \sum_{k=1}^3 \frac{\partial h_3}{\partial h_k} \cdot \frac{\partial h_k}{\partial W}$$


This iterative process is the essence of backpropagation through time.

Types Of Recurrent Neural Networks

There are four types of RNNs based on the number of inputs and outputs in the network:

1. One-to-One RNN


This is the simplest type of neural network architecture where there is a single input and a single output. It is used for straightforward classification tasks such as binary classification where no sequential data is involved.



One to One RNN

2. One-to-Many RNN


In a One-to-Many RNN the network processes a single input to produce multiple outputs over time. This is useful in tasks where one input triggers a sequence of predictions (outputs). For example in image captioning a single image can be used as input to generate a sequence of words as a caption.



One to Many RNN


3. Many-to-One RNN

The **Many-to-One RNN** receives a sequence of inputs and generates a single output. This type is useful when the overall context of the input sequence is needed to make one prediction. In sentiment analysis the model receives a sequence of words (like a sentence) and produces a single output like positive, negative or neutral.

*Many to One RNN*

4. Many-to-Many RNN

The **Many-to-Many RNN** type processes a sequence of inputs and generates a sequence of outputs. In language translation task a sequence of words in one language is given as input, and a corresponding sequence in another language is generated as output.




Generative Adversarial Networks (GANs)

Generative Adversarial Networks (GANs) were introduced by Ian Goodfellow and his colleagues in 2014. GANs are a class of neural networks that autonomously learn patterns in the input data to generate new examples resembling the original dataset.

GAN's architecture consists of two neural networks:

1. **Generator:** creates synthetic data from random noise to produce data so realistic that the discriminator cannot distinguish it from real data.
2. **Discriminator:** acts as a critic, evaluating whether the data it receives is real or fake.

They use adversarial training to produce artificial data that is identical to actual data.



The two networks engage in a continuous game of cat and mouse: the Generator improves its ability to create realistic data, while the Discriminator becomes better at detecting fakes. Over time, this adversarial process leads to the generation of highly realistic and high-quality data.

Detailed Architecture of GANs

The generator and discriminator model of GANs in detail:

1. Generator Model

The **generator** is a deep neural network that takes random noise as input to generate realistic data samples (e.g., images or text). It learns the underlying data distribution by adjusting its parameters through **backpropagation**.

The generator's objective is to produce samples that the discriminator classifies as real. The loss function is:

$$J_G = -\frac{1}{m} \sum_{i=1}^m \log D(G(z_i))$$

Where,

- J_G measure how well the generator is fooling the discriminator.
- $\log D(G(z_i))$ represents log probability of the discriminator being correct for generated samples.
- The generator aims to minimize this loss, encouraging the production of samples that the discriminator classifies as real ($\log D(G(z_i))$, close to 1).

2. Discriminator Model

The **discriminator** acts as a **binary classifier**, distinguishing between real and generated data. It learns to improve its classification ability through training, refining its parameters to **detect fake samples more accurately**.

When dealing with image data, the discriminator often employs **convolutional layers** or other relevant architectures suited to the data type. These layers help extract features and enhance the model's ability to differentiate between real and generated samples.

The discriminator reduces the negative log likelihood of correctly classifying both produced and real samples. This loss incentivizes the discriminator to accurately categorize generated samples as fake and real samples with the following equation:

$$J_D = -\frac{1}{m} \sum_{i=1}^m \log D(x_i) - \frac{1}{m} \sum_{i=1}^m \log(1 - D(G(z_i)))$$

- J_D assesses the discriminator's ability to discern between produced and actual samples.
- The log likelihood that the discriminator will accurately categorize real data is represented by $\log D(x_i)$.
- The log chance that the discriminator would correctly categorize generated samples as fake is represented by $\log(1 - D(G(z_i)))$.

By **minimizing this loss**, the discriminator becomes more effective at distinguishing between real and generated samples.

MinMax Loss


GANs follow a minimax optimization where the generator and discriminator are adversaries:

$$\min_G \max_D (G, D) = [\mathbb{E}_{x \sim p_{data}} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(g(z)))]]$$

Where,

- G is generator network and D is the discriminator network
- Actual data samples obtained from the true data distribution $p_{data}(x)$ are represented by x .
- Random noise sampled from a previous distribution $p_z(z)$ (usually a normal or uniform distribution) is represented by z .
- $D(x)$ represents the discriminator's likelihood of correctly identifying actual data as real.
- $D(G(z))$ is the likelihood that the discriminator will identify generated data coming from the generator as authentic.

The generator aims to **minimize** the loss, while the discriminator tries to **maximize** its classification accuracy.



How does a GAN work?

Let's understand how the generator (G) and discriminator (D) compete to improve each other over time:

1. Generator's First Move

G takes a random noise vector as input. This noise vector contains random values and acts as the starting point for G's creation process. Using its internal layers and learned patterns, G transforms the noise vector into a new data sample, like a generated image.

2. Discriminator's Turn

D receives two kinds of inputs:

- Real data samples from the training dataset.
- The data samples generated by G in the previous step.

D's job is to analyze each input and determine whether it's real data or something G cooked up. It outputs a probability score between 0 and 1. A score of 1 indicates the data is likely real, and 0 suggests it's fake.

3. Adversarial Learning

- If the discriminator correctly classifies real data as real and fake data as fake, it strengthens its ability slightly.
- If the generator successfully fools the discriminator, it receives a positive update, while the discriminator is penalized.

4. Generator's Improvement

Every time the discriminator misclassifies fake data as real, the generator learns and improves. Over multiple iterations, the generator produces more convincing synthetic samples.

5. Discriminator's Adaptation

The discriminator continuously refines its ability to distinguish real from fake data. This ongoing duel between the generator and discriminator enhances the overall model's learning process.

6. Training Progression

- As training continues, the generator becomes highly proficient at producing realistic data.
- Eventually, the discriminator struggles to distinguish real from fake, indicating that the GAN has reached a well-trained state.
- At this point, the generator can be used to generate high-quality synthetic data for various applications.

Types of GANs

1. Vanilla GAN

Vanilla GAN is the simplest type of GAN. It consists of:

- A generator and a discriminator, both are built using multi-layer perceptrons (MLPs).
- The model optimizes its mathematical formulation using stochastic gradient descent (SGD).

While Vanilla GANs serve as the foundation for more advanced GAN models, they often struggle with issues like mode collapse and unstable training.

2. Conditional GAN (CGAN)

Conditional GANs (CGANs) introduce an additional conditional parameter to guide the generation process. Instead of generating data randomly, CGANs allow the model to produce specific types of outputs.

Working of CGANs:

- A conditional variable (y) is fed into both the generator and the discriminator.
- This ensures that the generator creates data corresponding to the given condition (e.g., generating images of specific objects).
- The discriminator also receives the labels to help distinguish between real and fake data.

3. Deep Convolutional GAN (DCGAN)

Deep Convolutional GANs (DCGANs) are among the most popular and widely used types of GANs, particularly for image generation.

What Makes DCGAN Special?

- Uses Convolutional Neural Networks (CNNs) instead of simple multi-layer perceptrons (MLPs).
- Max pooling layers are replaced with convolutional stride, making the model more efficient.

- Fully connected layers are removed, allowing for better spatial understanding of images.

DCGANs have been highly successful in generating high-quality images, making them a go-to choice for deep learning researchers.

4. Laplacian Pyramid GAN (LAPGAN)

Laplacian Pyramid GAN (LAPGAN) is designed to generate ultra-high-quality images by leveraging a multi-resolution approach.

Working of LAPGAN:

- Uses multiple generator-discriminator pairs at different levels of the Laplacian pyramid.
- Images are first downsampled at each layer of the pyramid and upsampled again using Conditional GANs (CGANs).
- This process allows the image to gradually refine details, reducing noise and improving clarity.

Due to its ability to generate highly detailed images, LAPGAN is considered a superior approach for photorealistic image generation.

5. Super Resolution GAN (SRGAN)

Super-Resolution GAN (SRGAN) is specifically designed to increase the resolution of low-quality images while preserving details.

Working of SRGAN:

- Uses a deep neural network combined with an adversarial loss function.
- Enhances low-resolution images by adding finer details, making them appear sharper and more realistic.
- Helps reduce common image upscaling errors, such as blurriness and pixelation.

Autoencoders : Deeplearning model

An autoencoder is a type of **artificial neural network** that learns to represent data in a compressed form and then reconstructs it as closely as possible to the original input.

Autoencoders consists of two components:


- **Encoder:** This compresses the input into a compact representation and capture the most relevant features.
- **Decoder:** It reconstructs the input data from this compressed form to make it as similar as possible to the original input.

For example if the input is a noisy image of handwritten digits the autoencoder can learn to remove noise by compressing the image into a smaller feature set and reconstructing a cleaner version of the original image.

Autoencoders aim to minimize **reconstruction error** which is the difference between the input and the reconstructed output. They use loss functions such as **Mean Squared Error (MSE)** or **Binary Cross-Entropy (BCE)** and optimize through **backpropagation and gradient descent**. They are used in applications like image processing, anomaly detection, noise removal and feature extraction

Architecture of Autoencoder in Deep Learning

The architecture of an autoencoder consists of three main components: the Encoder, the Bottleneck (Latent Space) and the Decoder. Let's deep dive into each part to understand how they work together.



1. Encoder

The encoder is the part of the network that takes the input data and compresses it into a smaller lower-dimensional representation.

- **Input Layer:** This is where the original data enters the network e.g., an image or a set of features.
- **Hidden Layers:** These layers apply transformations to the input data. The encoder's goal is to extract essential features and reduce the data's dimensionality.
- **Output of Encoder (Latent Space):** The encoder outputs a compressed version of the data often called the latent representation or encoding. This is a condensed version of the input retaining only the important features.

2. Bottleneck (Latent Space)

The bottleneck is the smallest layer of the network where the data is represented in its most compressed form. It's often referred to as the latent space or code.

- This layer contains a reduced set of features representing the most important information from the input.
- The idea is that through this compression the network learns the key patterns and structures of the input data.

3. Decoder

The decoder is responsible for taking the compressed representation from the latent space and reconstructing it back into the original data form.

- **Hidden Layers:** The decoder uses a series of layers to gradually expand the compressed data back to the original input's dimensions.
- **Output Layer:** This layer produces the reconstructed data and aim to closely resemble the input data.

Loss Function in Autoencoder Training

During training an autoencoder aims to minimize the reconstruction loss which measures the difference between the original input and the reconstructed output. The choice of loss function depends on the type of data:

- **Mean Squared Error (MSE):** This is commonly used for continuous data. It measures the average squared differences between the input and the reconstructed data.
- **Binary Cross-Entropy:** Used for binary data (0 or 1 values). It calculates the difference in probability between the original and reconstructed output.

The network adjusts its weights to minimize this reconstruction loss learning to extract and retain only the most important features of the input data which are then stored in the latent space (bottleneck layer).

Autoencoder Efficient Representations

Constraining an autoencoders makes it learn and show efficient representation. Constraining an autoencoder means that the network learns meaningful, compact and useful features from the input data. After the network is trained only the encoder part is used to encode similar data for future tasks. Several techniques can be employed to achieve this goal:

- **Keep Small Hidden Layers:** By making each hidden layer as small as possible the network is forced to learn only the most important features of the data. A smaller layer size means less redundancy and efficient encoding.
- **Regularization:** Regularization methods like L1 or L2 regularization add a penalty terms to the loss function, preventing the network from overfitting. This ensures that the learned representations are more general and useful by discouraging excessively large weights.
- **Denoising:** In denoising autoencoders random noise is added to the input data during training. The network is then forced to remove this noise during reconstruction which helps it focus on the core, noise-free features of the data and makes the model more robust.
- **Tuning the Activation Functions:** Modifying the activation functions can encourage sparsity in the hidden layer, where only a few neurons are active at a time. This sparsity forces the network to learn only the most relevant features, reducing the complexity of the model and improving its efficiency.

Transformer Networks

Transformers are a type of deep learning model that utilizes **self-attention mechanisms** to process and generate sequences of data efficiently. They capture long-range dependencies and contextual relationships making them highly effective for tasks like language modeling, machine translation and text generation.

In this article we will explore the architecture and working of transformers by understanding their key components, mathematical formulations and how they function during training and inference.


Understanding Transformer Architecture

The transformer model is built on an encoder-decoder architecture where both the encoder and decoder are composed of a series of layers that utilize self-attention mechanisms and feed-forward neural networks. This architecture enables the model to process input data in parallel making it highly efficient and effective for tasks involving sequential data.

- The encoder processes input sequences and creates meaningful representations.

- The decoder generates outputs based on encoder representations and previously predicted tokens.

The encoder and decoder work together to transform the input into the desired output such as translating a sentence from one language to another or generating a response to a query.



Key Components Transformer

Transformers have 2 main components

1. Encoder

The primary function of the encoder is to create a high-dimensional representation of the input sequence that the decoder can use to generate the output. Encoder consists multiple layers and each layer is composed of two main sub-layers:

- Self-Attention Mechanism:** This sub-layer allows the encoder to weigh the importance of different parts of the input sequence differently to capture dependencies regardless of their distance within the sequence.
- Feed-Forward Neural Network:** This sub-layer consists of two linear transformations with a ReLU activation in between. It processes the output of the self-attention mechanism to generate a refined representation.

Layer normalization and residual connections are used around each of these sub-layers to ensure stability and improve convergence during training.

2. Decoder

Decoder in transformer also consists of multiple identical layers. Its primary function is to generate the output sequence based on the representations provided by the encoder and the previously generated tokens of the output.

Each decoder layer consists of three main sub-layers:

1. **Masked Self-Attention Mechanism:** Similar to the encoder's self-attention mechanism but with a mask to prevent the decoder from attending to future tokens in the output sequence.
2. **Encoder-Decoder Attention Mechanism:** This sub-layer allows the decoder to focus on relevant parts of the encoder's output representation, facilitating the generation of coherent and contextually appropriate output sequences.
3. **Feed-Forward Neural Network:** This sub-layer processes the combined output of the masked self-attention and encoder-decoder attention mechanisms.

How Transformers Work

1. Input Representation

The first step in processing input data involves converting raw text into a format that the transformer model can understand. This involves tokenization and embedding.

- **Tokenization:** The input text is split into smaller units called tokens, which can be words, subwords, or characters. Tokenization ensures that the text is broken down into manageable pieces.
- **Embedding:** Each token is then converted into a fixed-size vector using an embedding layer. This layer maps each token to a dense vector representation that captures its semantic meaning.
- Positional encodings are added to these embeddings to provide information about the token positions within the sequence.

2. Encoder Process in Transformers

1. **Input Embedding:** The input sequence is tokenized and converted into embeddings with positional encodings added.
2. **Self-Attention Mechanism:** Each token in the input sequence attends to every other token to capture dependencies and contextual information.
3. **Feed-Forward Network:** The output from the self-attention mechanism is passed through a position-wise feed-forward network.
4. **Layer Normalization and Residual Connections:** Layer normalization and residual connections are applied.

3. Decoder Process

1. **Input Embedding and Positional Encoding:** The partially generated output sequence is tokenized and embedded with positional encodings added.
2. **Masked Self-Attention Mechanism:** The decoder uses masked self-attention to prevent attending to future tokens ensuring that the model generates the sequence step-by-step.
3. **Encoder-Decoder Attention Mechanism:** The decoder attends to the encoder's output allowing it to focus on relevant parts of the input sequence.
4. **Feed-Forward Network:** Similar to the encoder the output from the attention mechanisms is passed through a position-wise feed-forward network.
5. **Layer Normalization and Residual Connections:** Similar to the encoder Layer normalization and residual connections are applied.

4. Training and Inference

Transformers are trained using supervised learning where the model learns to predict the next token in a sequence given the previous tokens.

Transformers have transformed deep learning by using self-attention mechanisms to efficiently process and generate sequences capturing long-range dependencies and contextual relationships. Their encoder-decoder architecture combined with multi-head attention and feed-forward networks enables highly effective handling of sequential data.

Deep Learning Applications

1. Computer vision

In computer vision, deep learning models enable machines to identify and understand visual data. Some of the main applications of deep learning in computer vision include:

- **Object detection and recognition:** Deep learning models are used to identify and locate objects within images and videos, making it possible for machines to perform tasks such as self-driving cars, surveillance, and robotics.
- **Image classification:** Deep learning models can be used to classify images into categories such as animals, plants, and buildings. This is used in applications such as medical imaging, quality control, and image retrieval.
- **Image segmentation:** Deep learning models can be used for image segmentation into different regions, making it possible to identify specific features within images.

2. Natural language processing (NLP)

In NLP, deep learning model enable machines to understand and generate human language. Some of the main applications of deep learning in NLP include:

- **Automatic Text Generation:** Deep learning model can learn the corpus of text and new text like summaries, essays can be automatically generated using these trained models.
- **Language translation:** Deep learning models can translate text from one language to another, making it possible to communicate with people from different linguistic backgrounds.
- **Sentiment analysis:** Deep learning models can analyze the sentiment of a piece of text, making it possible to determine whether the text is positive, negative, or neutral.
- **Speech recognition:** Deep learning models can recognize and transcribe spoken words, making it possible to perform tasks such as speech-to-text conversion, voice search, and voice-controlled devices.

3. Reinforcement learning

In reinforcement learning, deep learning works as training agents to take action in an environment to maximize a reward. Some of the main applications of deep learning in reinforcement learning include:

- **Game playing:** Deep reinforcement learning models have been able to beat human experts at games such as Go, Chess, and Atari.
- **Robotics:** Deep reinforcement learning models can be used to train robots to perform complex tasks such as grasping objects, navigation, and manipulation.
- **Control systems:** Deep reinforcement learning models can be used to control complex systems such as power grids, traffic management, and supply chain optimization.

Challenges in Deep Learning

Deep learning has made significant advancements in various fields, but there are still some challenges that need to be addressed. Here are some of the main challenges in deep learning:

1. **Data availability:** It requires large amounts of data to learn from. For using deep learning it's a big concern to gather as much data for training.
2. **Computational Resources:** For training the deep learning model, it is computationally expensive because it requires specialized hardware like GPUs and TPUs.

3. **Time-consuming:** While working on sequential data depending on the computational resource it can take very large even in days or months.
4. **Interpretability:** Deep learning models are complex, it works like a black box. it is very difficult to interpret the result.
5. **Overfitting:** when the model is trained again and again, it becomes too specialized for the training data, leading to overfitting and poor performance on new data.

Advantages of Deep Learning

1. **High accuracy:** Deep Learning algorithms can achieve state-of-the-art performance in various tasks, such as image recognition and natural language processing.
2. **Automated feature engineering:** Deep Learning algorithms can automatically discover and learn relevant features from data without the need for manual feature engineering.
3. **Scalability:** Deep Learning models can scale to handle large and complex datasets, and can learn from massive amounts of data.
4. **Flexibility:** Deep Learning models can be applied to a wide range of tasks and can handle various types of data, such as images, text, and speech.
5. **Continual improvement:** Deep Learning models can continually improve their performance as more data becomes available.

Disadvantages of Deep Learning

1. **High computational requirements:** Deep Learning AI models require large amounts of data and computational resources to train and optimize.
2. **Requires large amounts of labeled data:** Deep Learning models often require a large amount of labeled data for training, which can be expensive and time-consuming to acquire.
3. **Interpretability:** Deep Learning models can be challenging to interpret, making it difficult to understand how they make decisions.
Overfitting: Deep Learning models can sometimes overfit to the training data, resulting in poor performance on new and unseen data.
4. **Black-box nature:** Deep Learning models are often treated as black boxes, making it difficult to understand how they work and how they arrived at their predictions.

Difference between ANN, CNN and RNN

Artificial Neural Network (ANN):

Artificial Neural Network(ANN), is a group of multiple perceptrons or neurons at each layer. ANN is also known as a Feed-Forward Neural network because inputs are processed only in the forward direction. This type of neural networks are one of the simplest variants of neural networks. They pass information in one direction, through various input nodes, until it makes it to the output node. The network may or may not have hidden node layers, making their functioning more interpretable.

Advantages:

- Storing information on the entire network.
- Ability to work with incomplete knowledge.
- Having fault tolerance.
- Having a distributed memory.

Disadvantages:

- Hardware dependence.
- Unexplained behavior of the network.
- Determination of proper network structure.

Convolutional Neural Network (CNN):

Convolutional neural networks(CNN) are one of the most popular models used today. This neural network computational model uses a variation of multilayer perceptrons and contains one or more convolutional layers that can be either entirely connected or pooled. These convolutional layers create feature maps that record a region of image which is ultimately broken into rectangles and sent out for nonlinear processing.

Advantages:

- Very High accuracy in image recognition problems.
- Automatically detects the important features without any human supervision.
- Weight sharing.

Disadvantages:

- CNN do not encode the position and orientation of object.
- Lack of ability to be spatially invariant to the input data.
- Lots of training data is required.

Recurrent Neural Network (RNN):

Recurrent neural networks(RNN) are more complex. They save the output of processing nodes and feed the result back into the model (they did not pass the information in one direction only). This is how the model is said to learn to predict the outcome of a layer. Each node in the RNN model acts as a memory cell, continuing the computation and implementation of operations. If the network's prediction is incorrect, then the system self-learns and continues working towards the correct prediction during backpropagation.

Advantages:

- An RNN remembers each and every information through time. It is useful in time series prediction only because of the feature to remember previous inputs as well. This is called Long Short Term Memory.
- Recurrent neural network are even used with convolutional layers to extend the effective pixel neighborhood.

Disadvantages:

- Gradient vanishing and exploding problems.
- Training an RNN is a very difficult task.
- It cannot process very long sequences if using tanh or relu as an activation function.

Summation of all three networks in single table:

	<i>ANN</i>	<i>CNN</i>	<i>RNN</i>
Type of Data	Tabular Data, Text Data	Image Data	Sequence data

	<i>ANN</i>	<i>CNN</i>	<i>RNN</i>
Parameter Sharing	No	Yes	Yes
Fixed Length input	Yes	Yes	No
Recurrent Connections	No	No	Yes
Vanishing and Exploding Gradient	Yes	Yes	Yes
Spatial Relationship	No	Yes	No
Performance	ANN is considered to be less powerful than CNN, RNN.	CNN is considered to be more powerful than ANN, RNN.	RNN includes less feature compatibility when compared to CNN.
Application	Facial recognition and Computer vision.	Facial recognition, text digitization and Natural language processing.	Text-to-speech conversions.
Main advantages	Having fault tolerance, Ability to work with incomplete knowledge.	High accuracy in image recognition problems, Weight sharing.	Remembers each and every information, Time series prediction.

	<i>ANN</i>	<i>CNN</i>	<i>RNN</i>
Disadvantages	Hardware dependence, Unexplained behavior of the network.	Large training data needed, don't encode the position and orientation of object.	Gradient vanishing, exploding gradient.

Question Bank:

1. What is Neural Network? Explain its importance.
2. Explain working of Neural Network.
3. Enlist types of Neural Network. Explain its application.
4. Write advantages and disadvantages of Neural Network.
5. What is Deep learning? How deep learning works?
6. How is Deep Learning different from Machine Learning? Or compare Deep learning and Machine learning.
7. What is perceptron? Write its types.
8. Enlist deep learning models. Explain any two models in details.
9. Explain Feed forward neural networks (FNNs).
10. Explain Convolutional Neural Networks. (CNNs).
11. Explain Autoencoders.
12. Explain Transformer Networks.
13. Explain Recurrent Neural Network. (RNNs).
14. What is Deep Learning? Explain with examples. (3 Marks)
15. How is Deep Learning different from Machine Learning? (4 Marks)
16. Explain the key components of a Deep Learning model. (7 Marks)
17. What are Neural Networks? Define their basic structure. (3 Marks)
18. Discuss different types of activation functions in Neural Networks. (4 Marks)
19. Explain how backpropagation works in Neural Networks with a suitable example. (7 Marks)
20. List and define any three types of Deep Learning models. (3 Marks)
21. Compare CNN and RNN in terms of architecture and applications. (4 Marks)
22. Describe the architecture and working of Generative Adversarial Networks (GANs). (7 Marks)
23. Name three real-world applications of Deep Learning. (3 Marks)
24. How is Deep Learning applied in Natural Language Processing (NLP)? (4 Marks)
25. Discuss the impact of Deep Learning in the field of Computer Vision. (7 Marks)